

PicoCTF 2014 Writeups

Reed Koser

December 12, 2014

1 Police Records

In order to solve this problem, I needed to do two things: pull all of the badgeids from the picocTF server and then find the duplicate. This is a bit tricky because the badge server “encrypts” all communications with an xor key dictated by the server. Fortunately, both the encryption and the protocol is symmetric so all of the hard work of handling communication with the server has already been done for us in the provided source code.

Using the provided encryption/decryption routines, I wrote a short python client that iterated through badge IDs and printed them out. Once I knew that the badge IDs went from 0-1000 I extended the client to keep track of known badge ids and print out duplicates, solving the challenge.

2 Obfuscation

In order to solve obfuscation, the first step I took was to disassemble the binary. True to its name, it promised to be an exercise in frustration should I attempt to discern the function of the program statically. Instead, I used the linux `perf` tool to acquire information about the runtime of the program. The `perf stat` command can provide instruction-count resolution, and by maximizing the instruction count I was able to find an input which the binary accepted. It so happened that there was more than one key which the binary accepted and the one that my timing attack found first was not the key the game expected. However, a quick message to the mods and we had Obfuscation marked as solved.

3 Baleful

To solve Baleful, I began by running `strings` on the binary. The hint mentions that the binary is packed with a common packer, and the output of `strings` confirms this. Specifically, there was a string that says:

```
$ Info: This file is packed with the UPX executable packer http://upx.sf.net $
```

which gives away the name of the packer.

After unpacking the binary, I was confronted with thousands of machine code instructions. At this point a normal person would likely try to step through the program with a debugger, but after a few minutes of that I realized I wasn't going to get anywhere without understanding how the code was structured, so I rolled up my sleeves and wheeled out the nuclear weapon of binary analysis problems: IDA.

The first revelation came when I realized that the program was structured as a virtual machine. After this realisation I decided that the best solution would be to write a recursive descent disassembler (possibly not the most efficient solution, but source is available upon request!¹), in order to continue my quest through the land of machine code. Once I knew the general program flow, it was simply a matter of filling in opcodes, and in short order I had decoded the first section of code the virtual machine executes:

00001000	18 01 00 3a 10 00 00	mov r00 0000103a
00001007	18 01 01 db 1e 00 00	mov r01 00001edb
0000100e	18 00 03 00	mov r03 r00
00001012	18 01 05 42 cf 74 01	mov r05 0174cf42
00001019	1b 04 03	mov r04 [r03]
0000101c	06 00 04 04 05	xor r04 (r04 ^ r05)
00001021	1c 03 04	mov [r03] r04
00001024	02 01 03 03 04 00 00 00	add r03 (r03 + 00000004)
0000102c	17 00 01 03	test (r01 - r03)
00001030	14 19 10 00 00	jge 00001019
00001035	0e 3a 10 00 00	jmp 0000103a
0000103a	0e c0 1b 00 00	jmp 00001bc0

This is a relatively obvious xor “encryption” decoder, and clearly I need the normal version of the code. At this point I had figured out where in memory the instruction pointer and main memory were stored, so I set a conditional breakpoint in **gdb** for when the instruction pointer had advanced into the encrypted opcodes and then dumped the VM's main memory again. Once I had the full code, it was a relatively simple matter to figure out what the code was doing. The first thing that jumped out at me was the code like this:

00001142	0f 3f 10 00 00	call 0000103f
00001147	18 01 00 69 00 00 00	mov r00 00000069
0000114e	0f 3f 10 00 00	call 0000103f
00001153	18 01 00 6f 00 00 00	mov r00 0000006f
0000115a	0f 3f 10 00 00	call 0000103f
0000115f	18 01 00 6e 00 00 00	mov r00 0000006e
00001166	0f 3f 10 00 00	call 0000103f

Since the function at 0x103f is just a dispatch to a **ccall** that prints the character in **r00**, these code segments are what outputs to the user. After finding and labeling all of these output functions, it was a simple (if time consuming) matter to crunch backwards through the code to figure out what conditions caused the program to print “Congratulations!...” The most restrictive (and therefore interesting) seemed to be this loop:

000017d9	18 00 1e 02	mov r1e r02 # some sorta loop
000017dd	04 01 1e 1e 04 00 00 00	imul r1e (r1e * 00000004)
000017e5	18 00 00 0a	mov r00 r0a
000017e9	02 00 1e 1e 00	add r1e (r1e + r00)

¹I would include it here, but it makes the report four times longer without adding any substantive content

000017ee	18 00 03 1e	mov r03 r1e
000017f2	1b 04 03	mov r04 [r03]
000017f5	05 01 00 03 02 04 00 00	idiv r00 (r03 / 00000402)
000017fd	00	nop
000017fe	18 00 1e 03	mov r1e r03
00001802	04 01 1e 1e 04 00 00 00	imul r1e (r1e * 00000004)
0000180a	18 00 00 05	mov r00 r05
0000180e	02 00 1e 1e 00	add r1e (r1e + r00)
00001813	18 00 03 1e	mov r03 r1e
00001817	1b 03 03	mov r03 [r03]
0000181a	06 00 04 04 03	xor r04 (r04 ^ r03)
0000181f	18 00 1e 02	mov r1e r02
00001823	04 01 1e 1e 04 00 00 00	imul r1e (r1e * 00000004)
0000182b	18 00 00 09	mov r00 r09
0000182f	02 00 1e 1e 00	add r1e (r1e + r00)
00001834	18 00 03 1e	mov r03 r1e
00001838	1b 03 03	mov r03 [r03]
0000183b	18 00 1e 04	mov r1e r04
0000183f	18 00 00 03	mov r00 r03
00001843	17 00 1e 00	test (r1e - r00)
00001847	15 51 18 00 00	jne 00001851
0000184c	0e 58 18 00 00	jmp 00001858
00001851	18 01 01 01 00 00 00	mov r01 00000001
00001858	02 01 02 02 01 00 00 00	add r02 (r02 + 00000001)
00001860	18 00 03 01	mov r03 r01
00001864	18 00 1e 02	mov r1e r02
00001868	18 01 00 1e 00 00 00	mov r00 0000001e
0000186f	17 00 1e 00	test (r1e - r00)
00001873	11 d9 17 00 00	jlt 000017d9 # loop end

which (after staring at it for awhile) is another XOR decryption loop. The previous two hundred or so opcodes fill in two buffers, one of which this loop decrypts using the other as the key. Once I figured that out, I used `gdb` to dump the memory in those two buffers and a bit of python magic to xor the two and spit out the flag.